

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sb
5
6 from sklearn.model_selection import train_test_split
7 from sklearn.preprocessing import StandardScaler
8 from sklearn.linear_model import LogisticRegression
9 from sklearn.svm import SVC
10 from sklearn.preprocessing import MinMaxScaler
11 from xgboost import XGBClassifier
12 from sklearn import metrics
13
14 from keras.models import Sequential
15 from keras.layers import LSTM
16 from keras.layers import Dense
17 from keras.layers import Dropout
18
19 import warnings
20 warnings.filterwarnings('ignore')
```

```
1 from google.colab import drive
2 drive.mount('/content/drive')
3 df = pd.read_csv('drive/My Drive/AAPLmain.csv')
4 df.head()
```

Mounted at /content/drive

| | Date | Open | High | Low | Close | Adj Close | Volume |
|---|------------|----------|----------|----------|----------|-----------|-----------|
| 0 | 1980-12-12 | 0.128348 | 0.128906 | 0.128348 | 0.128348 | 0.099722 | 469033600 |
| 1 | 1980-12-15 | 0.122210 | 0.122210 | 0.121652 | 0.121652 | 0.094519 | 175884800 |
| 2 | 1980-12-16 | 0.113281 | 0.113281 | 0.112723 | 0.112723 | 0.087582 | 105728000 |
| 3 | 1980-12-17 | 0.115513 | 0.116071 | 0.115513 | 0.115513 | 0.089749 | 86441600 |
| 4 | 1980-12-18 | 0.118862 | 0.119420 | 0.118862 | 0.118862 | 0.092351 | 73449600 |

```
1 df.shape

(10664, 7)
```

```
1 df.describe()
```

| | Open | High | Low | Close | Adj Close | Volume |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 10664.000000 | 10664.000000 | 10664.000000 | 10664.000000 | 10664.000000 | 1.066400e+04 |
| mean | 17.214226 | 17.411780 | 17.021893 | 17.224528 | 16.533494 | 3.261870e+08 |
| std | 36.387008 | 36.832052 | 35.965057 | 36.417573 | 36.066055 | 3.374503e+08 |
| min | 0.049665 | 0.049665 | 0.049107 | 0.049107 | 0.038154 | 0.000000e+00 |
| 25% | 0.287946 | 0.296875 | 0.282366 | 0.289063 | 0.238623 | 1.201032e+08 |
| 50% | 0.491071 | 0.498884 | 0.483962 | 0.491071 | 0.407027 | 2.135056e+08 |
| 75% | 16.693303 | 16.851606 | 16.577144 | 16.695893 | 14.498148 | 4.058439e+08 |
| max | 182.630005 | 182.940002 | 179.119995 | 182.009995 | 180.683868 | 7.421641e+09 |

```
1 df.info()

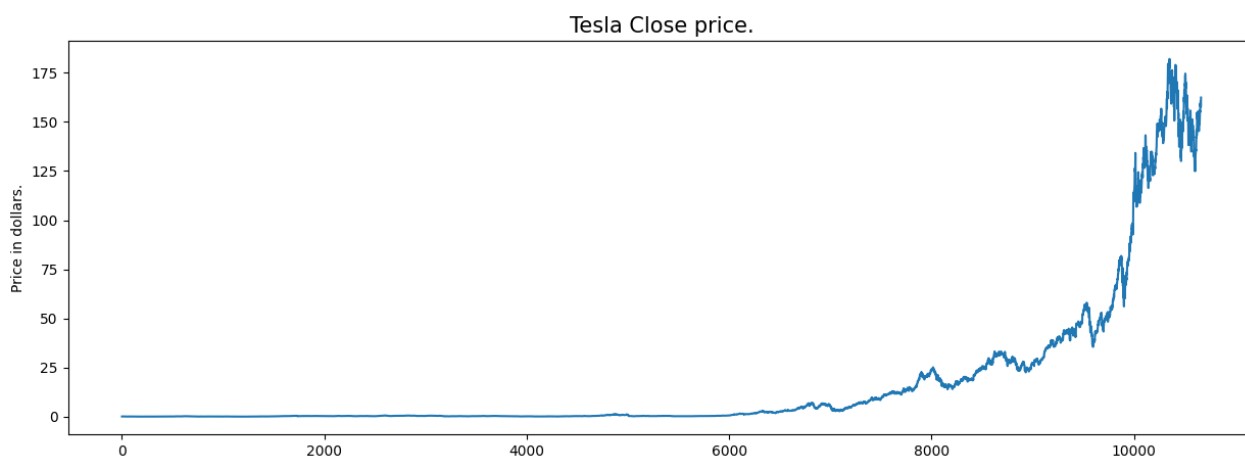
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10664 entries, 0 to 10663
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        10664 non-null  object
1   Open        10664 non-null  float64
2   High        10664 non-null  float64
3   Low         10664 non-null  float64
4   Close       10664 non-null  float64
5   Adj Close   10664 non-null  float64
6   Volume      10664 non-null  int64
dtypes: float64(5), int64(1), object(1)
memory usage: 583.3+ KB
```

```
1 plt.figure(figsize=(15,5))
2 plt.plot(df['Close'])
3 plt.title('AAPL Close price', fontsize=15)
```

```

3 plt.title('Tesla Close price.', fontsize=15)
4 plt.ylabel('Price in dollars.')
5 plt.show()

```



```
1 df.head()
```

| | Date | Open | High | Low | Close | Adj Close | Volume |
|---|------------|----------|----------|----------|----------|-----------|-----------|
| 0 | 1980-12-12 | 0.128348 | 0.128906 | 0.128348 | 0.128348 | 0.099722 | 469033600 |
| 1 | 1980-12-15 | 0.122210 | 0.122210 | 0.121652 | 0.121652 | 0.094519 | 175884800 |
| 2 | 1980-12-16 | 0.113281 | 0.113281 | 0.112723 | 0.112723 | 0.087582 | 105728000 |
| 3 | 1980-12-17 | 0.115513 | 0.116071 | 0.115513 | 0.115513 | 0.089749 | 86441600 |
| 4 | 1980-12-18 | 0.118862 | 0.119420 | 0.118862 | 0.118862 | 0.092351 | 73449600 |

```
1 df[df['Close'] == df['Adj Close']].shape
```

```
(34, 7)
```

```
1 df = df.drop(['Adj Close'], axis=1)
```

```
1 '''checking for null values'''
```

```
2 df.isnull().sum()
```

```

Date      0
Open      0
High      0
Low       0
Close     0
Volume    0
dtype: int64

```

```
1 features = ['Open', 'High', 'Low', 'Close', 'Volume']
```

```
2
```

```
3 plt.subplots(figsize=(20,10))
```


```
4
```

```
5 for i, col in enumerate(features):
```

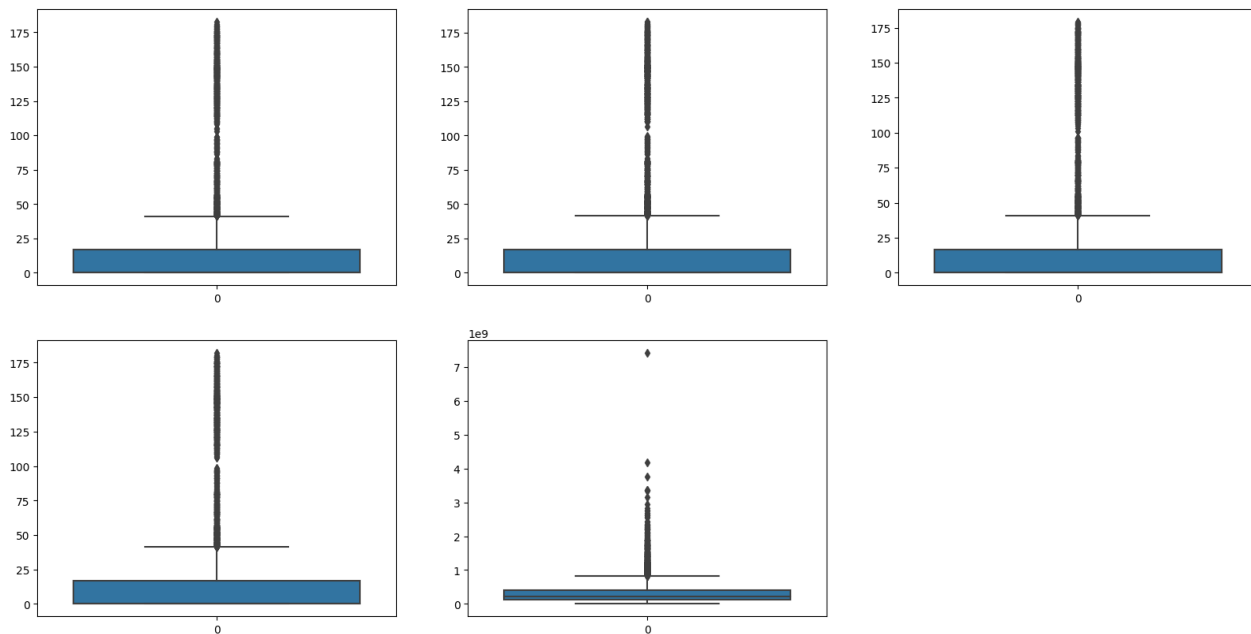
```
6     plt.subplot(2,3,i+1)
```

```
7     sb.distplot(df[col])
```

```
8 plt.show()
```



```
1 plt.subplots(figsize=(20,10))
2 for i, col in enumerate(features):
3     plt.subplot(2,3,i+1)
4     sb.boxplot(df[col])
5 plt.show()
```




```
1 splitted = df['Date'].str.split('-', expand=True)
2
3 df['day'] = splitted[2].astype('int')
4 df['month'] = splitted[1].astype('int')
5 df['year'] = splitted[0].astype('int')
6
7 df.head()
```

| | Date | Open | High | Low | Close | Volume | day | month | year |
|---|------------|----------|----------|----------|----------|-----------|-----|-------|------|
| 0 | 1980-12-12 | 0.128348 | 0.128906 | 0.128348 | 0.128348 | 469033600 | 12 | 12 | 1980 |
| 1 | 1980-12-15 | 0.122210 | 0.122210 | 0.121652 | 0.121652 | 175884800 | 15 | 12 | 1980 |
| 2 | 1980-12-16 | 0.113281 | 0.113281 | 0.112723 | 0.112723 | 105728000 | 16 | 12 | 1980 |
| 3 | 1980-12-17 | 0.115513 | 0.116071 | 0.115513 | 0.115513 | 86441600 | 17 | 12 | 1980 |
| 4 | 1980-12-18 | 0.118862 | 0.119420 | 0.118862 | 0.118862 | 73449600 | 18 | 12 | 1980 |

```
1 df['is_quarter_end'] = np.where(df['month']%3==0,1,0)
2 df.head()
```


| | Date | Open | High | Low | Close | Volume | day | month | year | is_quarter_end |
|---|------------|----------|----------|----------|----------|-----------|-----|-------|------|----------------|
| 0 | 1980-12-12 | 0.128348 | 0.128906 | 0.128348 | 0.128348 | 469033600 | 12 | 12 | 1980 | 1 |
| 1 | 1980-12-15 | 0.122210 | 0.122210 | 0.121652 | 0.121652 | 175884800 | 15 | 12 | 1980 | 1 |
| 2 | 1980-12-16 | 0.113281 | 0.113281 | 0.112723 | 0.112723 | 105728000 | 16 | 12 | 1980 | 1 |
| 3 | 1980-12-17 | 0.115513 | 0.116071 | 0.115513 | 0.115513 | 86441600 | 17 | 12 | 1980 | 1 |

```
1 data_grouped = df.groupby('day').mean()
2 plt.subplots(figsize=(20,10))
3
4 for i, col in enumerate(['Open', 'High', 'Low', 'Close']):
5     plt.subplot(2,2,i+1)
6     data_grouped[col].plot.bar()
7 plt.show()
```



```
1 df['open-close'] = df['Open'] - df['Close']
2 df['low-high'] = df['Low'] - df['High']
3 df['target'] = np.where(df['Close'].shift(-1) > df['Close'], 1, 0)

1 plt.pie(df['target'].value_counts().values,
2         labels=[0, 1], autopct='%1.1f%%')
3 plt.show()
```




```
1 plt.figure(figsize=(10, 10))
2
```

```

3 # As our concern is with the highly
4 # correlated features only so, we will visualize
5 # our heatmap as per that criteria only.
6 sb.heatmap(df.corr() > 0.9, annot=True, cbar=False)
7 plt.show()

```



```

1 features = df[['open-close', 'low-high', 'is_quarter_end']]
2 target = df['target']
3
4 scaler = StandardScaler()
5 features = scaler.fit_transform(features)
6
7 X_train, X_valid, Y_train, Y_valid = train_test_split(
8     features, target, test_size=0.1, random_state=2022)
9 print(X_train.shape, X_valid.shape)

```

(9597, 3) (1067, 3)


```

1 models = [LogisticRegression(), SVC(
2     kernel='poly', probability=True), XGBClassifier()]
3
4 for i in range(3):
5     models[i].fit(X_train, Y_train)
6
7     print(f'{models[i]} : ')
8     print('Training Accuracy : ', metrics.roc_auc_score(
9         Y_train, models[i].predict_proba(X_train)[:,-1]))
10    print('Validation Accuracy : ', metrics.roc_auc_score(
11        Y_valid, models[i].predict_proba(X_valid)[:,-1]))
12    print()

```

```
LogisticRegression() :  
Training Accuracy : 0.5269174389766803  
Validation Accuracy : 0.5309185412646317  
  
SVC(kernel='poly', probability=True) :  
Training Accuracy : 0.5174615996250385  
Validation Accuracy : 0.5082983533587819  
  
XGBClassifier(base_score=None, booster=None, callbacks=None,  
               colsample_bylevel=None, colsample_bynode=None,  
               colsample_bytree=None, early_stopping_rounds=None,  
               enable_categorical=False, eval_metric=None, feature_types=None,  
               gamma=None, gpu_id=None, grow_policy=None, importance_type=None,  
               interaction_constraints=None, learning_rate=None, max_bin=None,  
               max_cat_threshold=None, max_cat_to_onehot=None,  
               max_delta_step=None, max_depth=None, max_leaves=None,  
               min_child_weight=None, missing=nan, monotone_constraints=None,  
               n_estimators=100, n_jobs=None, num_parallel_tree=None,  
               predictor=None, random_state=None, ...) :  
Training Accuracy : 0.817039507331415  
Validation Accuracy : 0.5100080289606017
```

```
1 metrics.ConfusionMatrixDisplay.from_estimator(models[0], X_valid, Y_valid)  
2 plt.show()
```



```
1 dataset_train = pd.read_csv('drive/My Drive/AAPLmain.csv')  
2 dataset_train.head()
```

| | Date | Open | High | Low | Close | Adj Close | Volume |
|---|------------|----------|----------|----------|----------|-----------|-----------|
| 0 | 1980-12-12 | 0.128348 | 0.128906 | 0.128348 | 0.128348 | 0.099722 | 469033600 |
| 1 | 1980-12-15 | 0.122210 | 0.122210 | 0.121652 | 0.121652 | 0.094519 | 175884800 |
| 2 | 1980-12-16 | 0.113281 | 0.113281 | 0.112723 | 0.112723 | 0.087582 | 105728000 |
| 3 | 1980-12-17 | 0.115513 | 0.116071 | 0.115513 | 0.115513 | 0.089749 | 86441600 |
| 4 | 1980-12-18 | 0.118862 | 0.119420 | 0.118862 | 0.118862 | 0.092351 | 73449600 |

```
1 training_set = dataset_train.iloc[:,1:2].values  
2 print(training_set)  
3 print (training_set.shape)  
  
[[1.28348000e-01]  
 [1.22210000e-01]  
 [1.13281000e-01]  
 ...  
 [1.57970001e+02]  
 [1.59369995e+02]  
 [1.61529999e+02]]  
(10664, 1)  
  
1 scaler = MinMaxScaler (feature_range = (0,1))  
2 scaled_training_set = scaler.fit_transform(training_set)  
3 scaled_training_set
```

```
array([[4.30950014e-04],
       [3.97331936e-04],
       [3.48427437e-04],
       ...,
       [8.64936148e-01],
       [8.72603973e-01],
       [8.84434403e-01]])

1 X_train = []
2 y_train = []
3 for i in range(60,10664):
4     X_train.append(scaled_training_set[i-60:i, 0])
5     y_train.append(scaled_training_set[i, 0])
6 X_train = np.array(X_train)
7 y_train = np.array(y_train)

1 print( X_train.shape)
2 print( y_train.shape)

(10604, 60)
(10604,)

1 X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
2 X_train.shape

(10604, 60, 1)

1 regressor = Sequential()
2 regressor.add(LSTM(units = 50, return_sequences= True, input_shape = (X_train.shape[1], 1)))
3 regressor.add(Dropout (0.2))
4 regressor.add(LSTM(units = 50, return_sequences= True))
5 regressor.add(Dropout (0.2))
6 regressor.add(LSTM(units = 50, return_sequences=True))
7 regressor.add(Dropout (0.2))
8 regressor.add(LSTM(units = 50))
9 regressor.add (Dropout (0.2))
10 regressor.add(Dense (units=1))

1 regressor.compile(optimizer = 'adam', loss = 'mean_squared_error')
2 regressor.fit(X_train, y_train, epochs=100, batch_size=32)
```

```

332/332 [=====] - 34s 103ms/step - loss: 3.2040e-04
Epoch 94/100
332/332 [=====] - 34s 101ms/step - loss: 3.0982e-04
Epoch 95/100
332/332 [=====] - 33s 100ms/step - loss: 3.1802e-04
Epoch 96/100
332/332 [=====] - 32s 97ms/step - loss: 3.0795e-04
Epoch 97/100
332/332 [=====] - 33s 99ms/step - loss: 3.0633e-04
Epoch 98/100
332/332 [=====] - 33s 100ms/step - loss: 2.7265e-04
Epoch 99/100
332/332 [=====] - 34s 102ms/step - loss: 3.2877e-04
Epoch 100/100
332/332 [=====] - 34s 102ms/step - loss: 3.2680e-04
<keras.callbacks.History at 0x7f3235e6620>

```

```

1 dataset_test = pd.read_csv("drive/My Drive/TSLA2023-Jan-March.csv")
2 actual_stock_price = dataset_test.iloc[:,1:2].values

```

```

1 dataset_total = pd.concat((dataset_train [ 'Open'], dataset_test['Open']), axis = 0)
2 inputs = dataset_total[len(dataset_total)- len(dataset_test)-60:].values
3
4 inputs = inputs.reshape(-1,1)
5 inputs = scaler.transform(inputs)
6 inputs.shape
7
8 X_test = []
9 for i in range(60,120):
10     X_test.append(inputs [i-60:i, 0])
11 X_test = np.array(X_test)
12 X_test = np.reshape(X_test, (X_test.shape[0],X_test.shape[1] , 1))

```

```

1 predicted_stock_price = regressor.predict(X_test)
2 predicted_stock_price = scaler.inverse_transform(predicted_stock_price)
3

```

```

2/2 [=====] - 1s 28ms/step

```

```


1 plt.plot(actual_stock_price, color = 'red', label = 'Actual Apple Stock Price')
2 plt.plot(predicted_stock_price, color = 'blue', label = 'Predicted Apple Stock Price')
3 plt.title('Apple Stock Price Prediction')
4 plt.xlabel('Time')
5 plt.ylabel('Apple Stock Price')
6 plt.legend()

```

```

<matplotlib.legend.Legend at 0x7f32365063e0>

```



✓ 0s completed at 11:08 AM

● ×